

UK Crime Data Analysis with Apache Spark SQL

Nisha Khadka BK

MSc Data Analytics, De Montfort University

Big Data Analytics Applications (CSIP5203), Task 1

Task 1: Analysing Data with Spark SQL

1.1 Introduction of dataset

The dataset used in this analysis is acquired from UK Police street crime data which is publicly available at data.police.uk. It is made up of monthly police crime statistics from four UK counties:

- Essex
- Leicestershire
- Nottinghamshire
- Warwickshire

It spans over November 2023 to November 2025. According to the UK Open Government Licence framework, this data is made freely available to support research, transparency, and evidence-based policy development (UK Government, 2025). The selection of this particular dataset aligns with recommendations by Provost and Fawcett (2013), who emphasize that effective data science projects should work with real-world datasets that demonstrate both sufficient scale and practical relevance.

```
# List all month folders
month_folders = sorted([f for f in os.listdir(
    if os.path.isdir(os.
print(f"\n Found {len(month_folders)} mon
for i, folder in enumerate(month_folders, 1)
    print(f" {i}. {folder}")
```

Found 25 month folders:

1. 2023-11
2. 2023-12
3. 2024-01
4. 2024-02
5. 2024-03
6. 2024-04
7. 2024-05
8. 2024-06
9. 2024-07
10. 2024-08
11. 2024-09
12. 2024-10
13. 2024-11
14. 2024-12
15. 2025-01
16. 2025-02
17. 2025-03
18. 2025-04
19. 2025-05
20. 2025-06
21. 2025-07
22. 2025-08
23. 2025-09
24. 2025-10
25. 2025-11

```
# Defining Counties to analyse
counties = ['essex', 'leicestershire', 'nottinghamshire', 'warwickshire']
print(f"Counties to load: {'', '.join(counties)}")
```

Counties to load: essex, leicestershire, nottinghamshire, warwickshire

Figure 1: Data sets

Key details of our dataset:

- As stated by Zaharia et al. (2016) in their foundational work on Spark architecture, the dataset offers a significant volume to illustrate the capabilities of Apache Spark and distributed computing frameworks, with a total of 912,621 crime incidents recorded across the four counties over the 25-month period.
- Data structure: As outlined in the UK Police data schema documentation (data.police.uk, 2025), these attributes include:
 - ❖ Each crime has unique identified crime ID.
 - ❖ Report's month and year (YYYY-MM)
 - ❖ Police unit that documented the crime reported it.
 - ❖ Falls within territory of the police force where crime was committed
 - ❖ Location: Street name or location description; Longitude, Latitude: crime's geographic coordinates
 - ❖ Lower layer super output area (LSOA code, LSOA name)
 - ❖ Crime type: kind of crime (such as stealing or assault)

During loading, following metadata were added for temporal and categorical analysis:

- ❖ Year, month, year_month
- ❖ Date (first day of month), county, month_name

To enhance the analytical capabilities of the dataset, additional metadata columns were engineered during the data loading process, following best practices recommended by McKinney (2022) for temporal data manipulation in pandas and distributed computing frameworks.

1.2 Data Loading and Validation:

Data was imported from 100 csv files: 25 months of 4 counties and process was implemented using Apache Spark, leveraging PySpark's Dataframe API. As Karau et al. (2015) explain in their comprehensive guide to Spark, the framework's lazy evaluation and distributed computing paradigm ensure optimal performance even with limited computational resources. Zaharia et al. (2016) further demonstrate that Spark's in-memory processing capabilities can outperform conventional MapReduce frameworks for iterative algorithms by up to 100 times. The loading technique follows the multi-step procedure specified in an official PySpark MLlib manual (Anthropic, 2025):

- Each month's data was kept in a different subdirectory. YYYY-MM-county-street.csv was used to name each county's data for a certain month confirming to the UK Police data distribution format (data.police.uk, 2025).
- Instead of attempting to load all files simultaneously, a sequential loading approach was employed where each CSV file was read individually. McKinney (2022) notes that this incremental approach prevents memory overflow and provides progress feedback during execution, which is particularly valuable for debugging large-scale data pipelines.
- Record count verify 912,621 total records, schema inspection to guarantee correct column types and null value evaluation to comprehend data quality. VanderPlas

(2016) emphasizes that such validation steps are critical components of any robust data science workflow.

```
# Cache DataFrame for better performance

total_records = crime_data.count()

print("\n LOADING COMPLETE!")
print(f"Time period:{month_folders[0]} to {month_folders[-1]}")
print(f"Counties: {' '.join(counties)}")
print(f"Total files: {len(allfile_paths)}")
print(f"Total records: {total_records:,}")
```

LOADING POLICE DATA
=====

Loading files...
Found 100 CSV files.

Loading files.
Loaded 10 files
Loaded 20 files
Loaded 30 files
Loaded 40 files
Loaded 50 files
Loaded 60 files
Loaded 70 files
Loaded 80 files
Loaded 90 files
Loaded 100 files

All files loaded: 100 files

[Stage 201:=====> (80 + 4) / 100]

LOADING COMPLETE!
Time period:2023-11 to 2025-11
Counties: essex, leicestershire, nottinghamshire, warwickshire
Total files: 100
Total records: 912,621

Figure 2: Data Loading

```

DATA VALIDATION
=====

Total records: 912,621

Sample data:
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| Crime ID | Month | Reported by | Falls within | Longitude | Latitude | Location | LSOA code | LSOA name | Crime type | Last outcome category | Context | county | month_folder |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 4b4d60625323ffcc0... | 2023-11-01 00:00:00 | Essex Police | Essex Police | 0.709247 | 52.022613 | On or near Rogers... | E010029081 | Savergh 0020 | Other theft | Status update un... | NULL | Essex | 2023-11 |
| 49e904757038263f5... | 2023-11-01 00:00:00 | Essex Police | Essex Police | 0.432099 | 51.640106 | On or near Oakwo... | E010021237 | Basildon 001A | Violence and sexu... | Local resolution | NULL | Essex | 2023-11 |
| 26a67b541247fc121... | 2023-11-01 00:00:00 | Essex Police | Essex Police | 0.434079 | 51.641863 | On or near Longtail... | E010021238 | Basildon 001B | Violence and sexu... | Unable to prosecu... | NULL | Essex | 2023-11 |
| 22110400997b2546a... | 2023-11-01 00:00:00 | Essex Police | Essex Police | 0.432022 | 51.644629 | On or near Brooks... | E010021238 | Basildon 001B | Violence and sexu... | Unable to prosecu... | NULL | Essex | 2023-11 |
| 3f100001212f5f5a4... | 2023-11-01 00:00:00 | Essex Police | Essex Police | 0.426594 | 51.640023 | On or near Double... | E010021238 | Basildon 001B | Other crime | Unable to prosecu... | NULL | Essex | 2023-11 |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|

only showing top 5 rows

All columns:
- Crime ID
- Month
- Reported by
- Falls within
- Longitude
- Latitude
- Location
- LSOA code
- LSOA name
- Crime type
- Last outcome category
- Context
- county
- month_folder

Quick check:
has 'county' column: Yes
has 'year' column: No

=====
done!

```

Figure 3: Data validation

1.3 Data cleaning

```

=====
Data Cleaning and Preprocessing
=====

Missing values in critical columns:

Crime type      :      0 missing_data (  0.0%)

Longitude       : 14,816 missing_data (  1.6%)

Latitude        : 14,816 missing_data (  1.6%)

Location        :      0 missing_data (  0.0%)

[Stage 393:=====> (93 + 4) / 100]
county          :      0 missing_data (  0.0%)

Check complete.

```

Figure 4: Missing data

1.4 Data Analysis

SparkSQL was used to run a variety of analytical queries. Analyses were designed to address multiple dimensions, including geographical distribution, temporal patterns, and crime type prevalence. This multi-dimensional analytical approach aligns with recommendations by Provost and Fawcett (2013), who mentioned that comprehensive business intelligence requires examining data from various perspectives.

A. Total crimes distribution by county

Findings and Interpretation:

County	Crime Count
Essex	330,416(36.2%)
Nottinghamshire	254,860(27.9%)
Leicestershire	223,259(24.5%)
Warwickshire	104,086(11.4%)

The distribution suggests various factors: differences in population density, urbanization levels. Essex's high number may be related to its higher population or different methods of reporting crimes. Additionally, as Wickham and Grolemond (2017) caution in their work on data visualization and interpretation, regional disparities may be caused by changes in reporting procedures, demographic makeup, and economic circumstances.

Visualization:




Figure 5: Crime incidents by county

B. Most common crime types

Findings and Interpretation:

This analysis gave distribution of crime type and calculated their percentage contributions to overall crime landscape, following analytical frameworks outlined by VanderPlas (2016) for categorical data analysis.

- Violence and sexual offences turned out as the most common category representing 326,550 incidents of crime.
- Anti-social behavior ranked second with 124,236 incidents.

Domination of Violence and sexual offences draws attention to public safety that need different tactics. As McKinney (2022) discusses in context of data-driven public policy, understanding crime type distributions is essential for allocating investigative resources and developing specialized training programs.




Figure 6: Distribution of various crime

C. Year-over-year comparison of crimes

Findings and Interpretation:

This trend shows massive, one-time surge from 2023 to 2024 followed by a sensible decline across all counties in 2025.

- The increase from 2023 to 2024 are not considered normal variations. This suggests significant change taking place at start of 2024.
- The crime incidents decrease in every county ranging from -3.1% to -9.6%. It's safe to say the current situation is as steady with early indications of decline.

According to VanderPlas (2016), even little year-over-year variations should be evaluated cautiously. Potential contributing factors may include economic pressures, population growth, or changes in crime reporting practices, all of which require further investigation as recommended in the UK Government's evidence-based policing framework (UK Government, 2025).

Visualization:

ANALYSIS 3: Year-over-Year Crime Trends

Creating 'year' and 'month' columns from 'Month'...
Total Crimes by Year:

year	total_crimes
2023	70009
2024	439165
2025	403447

Yearly Crime Trends by County:

year	county	crime_count
2023	Essex	26402
2024	Essex	158403
2025	Essex	145611
2023	Leicestershire	16925
2024	Leicestershire	107902
2025	Leicestershire	98432
2023	Nottinghamshire	19178
2024	Nottinghamshire	123798
2025	Nottinghamshire	111884
2023	Warwickshire	7504
2024	Warwickshire	49062
2025	Warwickshire	47520

Pivot View (County × Year):

year	2023	2024	2025
county			
Essex	26402.0	158403.0	145611.0
Leicestershire	16925.0	107902.0	98432.0
Nottinghamshire	19178.0	123798.0	111884.0
Warwickshire	7504.0	49062.0	47520.0

Year-over-Year Growth Analysis:

Essex:

2023 → 2024: 26,402 → 158,403
Change: +500.0% (Significant Increase)
2024 → 2025: 158,403 → 145,611
Change: -8.1% (Decrease)

Leicestershire:

2023 → 2024: 16,925 → 107,902
Change: +537.5% (Significant Increase)
2024 → 2025: 107,902 → 98,432
Change: -8.8% (Decrease)




Figure 7: Distribution of crimes across years.

D. Crime Trend by county over years

Interpretation:

Looking at how crime has changed over years, we can get clear picture of each month. We tracked total number of incidents reported each month and searched for any pattern or trend.

- About 36,505 crimes has been reported per month.
- In July 2024, there were 42,187 incidents.
- February 2024 had 31,049 incidents making it quietest month.

Visualization:




Figure 8: Crime trends by county over years.

E. Seasonal crime distribution by County

Findings and Interpretation:

We can clearly see that most crime in all four counties is reported mostly in Autumn.

- Essex: 91,199 crimes
- Nottinghamshire: 69,439 crimes
- Leicestershire: 61,496 crimes
- Warwickshire: 28,888 crimes

Visualization:




Figure 9: Seasonal crime distribution by county

F. Top 10 Areas Crime Distribution

Findings and Interpretation:

- About 40% of total crimes occurs in only two locations: Leicester 040F: 22.6% and Leicester 041D: 17.8%.
- Rest of the distribution shows only 6 to 11 % in other eight locations.
This can imply that: Leicester has the biggest impact if policing and crime prevention efforts are shown.

Visualization:




Figure 10: Crime distribution in top areas.

G. Geographic Crime Distribution by County

Findings and Interpretation:

This analysis utilized Spark SQL window functions with PARTITION BY clauses to rank crime types independently for each geographical area, a technique thoroughly explained in the Spark documentation (Zaharia et al., 2016; Anthropic, 2025).

- Essex: Crime is fixated in the southeast area.
- Leicestershire: Crime clusters intensely around Leicester city.
- Nottinghamshire: Crime is heavily focused near Nottingham.
- Warwickshire: Crime centers on the Coventry/Nuneaton area in north.


Visualization:



```
Plotting crime locations for Leicestershire...
```

```
Found 216,745 crimes with coordinates
```


```
Sampling 9,761 points for visualization
```



```
Plotting crime locations for Nottinghamshire...
```

```
Found 252,932 crimes with coordinates
```

```
Sampling 9,915 points for visualization
```



Plotting crime locations for Warwickshire...

Found 103,132 crimes with coordinates

Sampling 10,148 points for visualization




Figure 11: Geographic Crime Distribution by county

Discussion and Findings

The analysis of UK police crime data covers 912,621 crime incidents across four counties over 25 months. The results of this analysis may inform appropriate choices of police resource allocation and law enforcement policies. The application of predictive analytics in this study demonstrates how unprocessed data may be converted into insightful and useful knowledge.

Using Spark's dataframes and SQL, we were able to compute for large-scale data. We could handle 900,000+ records and identify regional trends (autumn crime rises) since Spark could process the entire dataset at once, as recommended by Karau et al. (2015). Complex, non-linear interactions between characteristics that are challenging to identify using conventional statistical techniques can be found using machine learning analytics (Han, Kamber & Pei, 2012).

This report implies that policing resources, training programs, may need to be allocated evenly to address the differences. However, other crimes must be evaluated with caution. This suggests that present strategies and tactics are effective at maintaining relative stability. Main lesson emerged was appropriate visualization techniques significantly enhanced insight communication, making complex patterns accessible to non-technical stakeholders, a principle strongly advocated by Wickham and Grolemund (2017).

References and Bibliography

1. Brownlee, J. (2020) Machine learning mastery with Python: understand your data, create accurate models and work projects end-to-end. Vermont: Machine Learning Mastery.
2. data.police.uk (2025) UK police data portal: street-level crime data. Available at: <https://data.police.uk/data/> (Accessed: 28 January 2026).
3. Géron, A. (2019) Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow: concepts, tools, and techniques to build intelligent systems. 2nd edn. Sebastopol, CA: O'Reilly Media.
4. Han, J., Kamber, M. and Pei, J. (2012) Data mining: concepts and techniques. 3rd edn. Waltham, MA: Morgan Kaufmann.
5. Kaggle (2025) Titanic: machine learning from disaster. Available at: <https://www.kaggle.com/c/titanic> (Accessed: 28 January 2026).
6. Karau, H., Konwinski, A., Wendell, P. and Zaharia, M. (2015) Learning Spark: lightning-fast big data analysis. Sebastopol, CA: O'Reilly Media.
7. McKinney, W. (2022) Python for data analysis: data wrangling with pandas, NumPy, and Jupyter. 3rd edn. Sebastopol, CA: O'Reilly Media.
8. Provost, F. and Fawcett, T. (2013) Data science for business: what you need to know about data mining and data-analytic thinking. Sebastopol, CA: O'Reilly Media.
9. Raschka, S. and Mirjalili, V. (2019) Python machine learning: machine learning and deep learning with Python, scikit-learn, and TensorFlow 2. 3rd edn. Birmingham: Packt Publishing.
10. Spark (2025) PySpark MLlib documentation: machine learning library guide. Available at: <https://spark.apache.org/docs/latest/ml-guide.html> (Accessed: 28 January 2026).
11. UK Government (2025) Open government licence for public sector information. Available at: <http://www.nationalarchives.gov.uk/doc/open-government-licence/version/3/> (Accessed: 28 January 2026).
12. VanderPlas, J. (2016) Python data science handbook: essential tools for working with data. Sebastopol, CA: O'Reilly Media.
13. Wickham, H. and Grolemund, G. (2017) R for data science: import, tidy, transform, visualize, and model data. Sebastopol, CA: O'Reilly Media.
14. Zaharia, M., Xin, R.S., Wendell, P., Das, T., Armbrust, M., Dave, A., Meng, X., Rosen, J., Venkataraman, S., Franklin, M.J., Ghodsi, A., Gonzalez, J., Shenker, S. and Stoica, I. (2016) 'Apache Spark: a unified engine for big data processing', Communications of the ACM, 59(11), pp. 56-65. doi: 10.1145/2934664.

Other Resources

1. Google Colab Platform: Computational platform for running PySpark and Jupyter notebooks without local installations requirements.
2. Apache Spark Documentation: Official PySpark API references and example used for implementing Spark SQL queries and Mlib functions.
3. Matplotlib Documentation: Visualization libraries used for creating charts and graphs.
4. Demontfort Learning Zone: course materials, lecture slides and lab workshop.